

Multimodal Video Analysis Application

Multimodal Video Analysis Application

1. Concept Introduction
 - 1.1 What is "Video Analysis"?
 - 1.2 Implementation Principle Overview
2. Code Analysis
 - Key Code
 1. Tool Layer Entry (`large_model/utils/tools_manager.py`)
 2. Model Interface Layer and Frame Extraction (`large_model/utils/large_model_interface.py`)
 - Code Analysis
3. Practical Operations
 - 3.1 Configuring the Offline Large Model
 - 3.1.1 Configuring the LLM Platform (`yahboom.yaml`)
 - 3.1.2 Configuring the Model Interface (`large_model_interface.yaml`)
 - 3.2 Starting and Testing (Text Input Mode)
4. Common Problems and Solutions
 - Issue 1: Error message "Cannot find video file" or "Unable to extract keyframes from video."
 - Issue 2: Analyzing a long video is very time-consuming.

Note: Due to performance limitations, the Raspberry Pi 5 models with 1GB, 2GB, and 4GB of memory cannot run offline. You can refer to the tutorials for online large models.

1. Concept Introduction

1.1 What is "Video Analysis"?

In the `large_model` project, the **multimodal video analysis** feature enables a robot to process a video and summarize the core content of the video, describe key events, or answer specific questions about the video in natural language. This enables the robot to leap from understanding only static images to understanding dynamic and temporal relationships.

The core tool for this feature is `analyze_video`. When a user provides a video file and asks a question (such as "Summarize what this video says"), the system invokes this tool to process and analyze the video and return a textual response from the AI.

1.2 Implementation Principle Overview

The core challenge of offline video analysis is how to efficiently process video data containing hundreds or thousands of frames. A popular implementation principle is as follows:

1. **Keyframe Extraction:** First, the system does not process every frame of the video. Instead, it extracts the most representative keyframes through algorithms (such as scene boundary detection or fixed-time sampling). This significantly reduces the amount of data required for processing.
2. **Image Encoding:** Each extracted keyframe is fed into a visual encoder, similar to visual understanding applications, and converted into a digital vector containing image information.

3. **Temporal Information Fusion:** This is the most significant difference from single-image understanding. The model needs to understand the temporal order between these keyframes. Typically, a recurrent neural network (RNN) or Transformer model is used to fuse the vectors of all keyframes to form a "memory" vector that represents the dynamic content of the entire video.
4. **Question-Answering and Generation:** The user's text question is encoded and then cross-modally fused with this "memory" vector. Finally, the language model generates a summary of the entire video or an answer to a specific question based on this fused information.

Simply put, **it condenses a video into a few key images and their sequence, allowing users to understand the entire story like reading a comic strip and answer related questions.**

2. Code Analysis

Key Code

1. Tool Layer Entry (`largemodel/utlils/tools_manager.py`)

The `analyze_video` function in this file defines the tool's execution flow.

```
# From largemodel/utlils/tools_manager.py
class ToolsManager:
    # ...
    def analyze_video(self, args):
        """
        Analyze video file and provide content description.

        :param args: Arguments containing video path.
        :return: Dictionary with video description and path.
        """
        self.node.get_logger().info(f"Executing analyze_video() tool with args:
{args}")
        try:
            video_path = args.get("video_path")
            # ... (Intelligent path fallback mechanism)

            if video_path and os.path.exists(video_path):
                # ... (Building Prompt)

                # Use a fully isolated, one-time context for video analysis to
                ensure a plain text description.
                simple_context = [{
                    "role": "system",
                    "content": "You are a video description assistant. ..."
                }]

                result = self.node.model_client.infer_with_video(video_path,
prompt, message=simple_context)

                # ... (Processing results)
                return {
                    "description": description,
                    "video_path": video_path
```

```
}  
# ... (Error Handling)
```

2. Model Interface Layer and Frame Extraction (`largemodel/utils/large_model_interface.py`)

The functions in this file are responsible for processing video files and passing them to the underlying model.

```
# From largemodel/utils/large_model_interface.py  
  
class model_interface:  
    # ...  
    def infer_with_video(self, video_path, text=None, message=None):  
        """Unified video inference interface."""  
        # ... (Prepare Message)  
        try:  
            # Determine which specific implementation to call based on  
            self.llm_platform  
            if self.llm_platform == 'ollama':  
                response_content = self.ollama_infer(self.messages,  
video_path=video_path)  
            # ... (The logic of other online platforms)  
        # ...  
        return {'response': response_content, 'messages': self.messages.copy()}  
  
    def _extract_video_frames(self, video_path, max_frames=5):  
        """Extract keyframes from a video for analysis. """  
        try:  
            import cv2  
            # ... (Video reading and frame interval calculation)  
            while extracted_count < max_frames:  
                # ... (Looping through video frames)  
                if frame_count % frame_interval == 0:  
                    # ... (Save the frame as a temporary image)  
                    frame_base64 = self.encode_file_to_base64(temp_path)  
                    frame_images.append(frame_base64)  
                # ...  
            return frame_images  
        # ... (Exception handling)
```

Code Analysis

The implementation of video analysis is more complex than image analysis. It requires a key preprocessing step at the model interface layer: frame extraction.

1. Tools Layer (`tools_manager.py`):

- The `analyze_video` function is the entry point for video analysis. Its responsibilities are clear: accept a video file path and construct a prompt to request a description of the video content.

- It initiates the analysis process by calling the `self.node.model_client.infer_with_video` method. Like the visual understanding tool, it is completely agnostic about the underlying model details and only handles passing the "video file" and "analysis instructions."

2. Model Interface Layer (`large_model_interface.py`):

- The `infer_with_video` function is the scheduler that connects the upper-level tools with the underlying model. It dispatches tasks to the corresponding implementation functions based on the platform configuration (`self.llm_platform`).
- Unlike processing single images, processing videos requires an additional step. The `_extract_video_frames` method demonstrates the general logic for implementing this step: it uses the `cv2` library to read the video and extract several keyframes (by default, five).
- Each extracted frame is treated as a separate image and typically encoded as a Base64 string.
- Ultimately, a request containing multiple frame image data and analysis instructions is sent to the main model. The model performs a comprehensive analysis of these consecutive images to generate a description of the entire video content.
- This "video-to-multiple-images" preprocessing is completely encapsulated in the model interface layer and is transparent to the tool layer.

In summary, the general process of video analysis is: `ToolsManager` initiates an analysis request -> `model_interface` intercepts the request and calls `_extract_video_frames` to decompose the video file into multiple keyframe images -> `model_interface` sends these images, along with analysis instructions, to the corresponding model platform according to the configuration -> the model returns a comprehensive description of the video -> the results are finally returned to `ToolsManager`. This design ensures the stability and versatility of upper-layer applications.

3. Practical Operations

3.1 Configuring the Offline Large Model

3.1.1 Configuring the LLM Platform (`yahboom.yaml`)

This file determines which large model platform the `model_service` node loads as its primary language model.

First, enter the following command in the terminal to enter the Docker container:

```
./ros2_docker.sh
```

If you need to enter the same Docker container later to run other commands, simply enter `./ros2_docker.sh` again in the host terminal.

1. Open the file in the terminal:

```
vim ~/yahboom_ws/src/largemodel/config/yahboom.yaml
```

2. Modify/Confirm `llm_platform`:

```

model_service:                                #Model server node parameters
  ros__parameters:
    language: 'zh'                            #Large Model Interface Language
    useolinetts: True                         #This item is invalid in text mode
    and can be ignored

    # Large model configuration
    llm_platform: 'ollama'                    # Key: Make sure it's 'ollama'
    regional_setting : "china"
    camera_type: 'csi'                        # Camera type: 'csi', 'usb'
    mjpeg_stream_url: 'http://172.17.0.1:8080/camera.mjpg' # MJPEG stream URL
    for CSI camera in Docker

```

3.1.2 Configuring the Model Interface (large_model_interface.yaml)

This file defines which visual model to use when the `ollama` platform is selected.

1. Open the file in a terminal.

```
vim ~/yahboom_ws/src/largemodel/config/large_model_interface.yaml
```

2. Find the Ollama configuration

```

#.....
## Offline Large Language Models
# Ollama Configuration
ollama_host: "http://localhost:11434" # ollama server address
ollama_model: "llava" # Key: Replace this with the multimodal model you
downloaded, such as "llava"
#.....

```

Note: Please ensure that the model specified in the configuration parameters (e.g., `llava`) can handle multimodal input.

3.2 Starting and Testing (Text Input Mode)

Note: Using a model with a small number of parameters or running at a low memory usage will result in poor performance. For a better experience with this feature, please refer to the corresponding section in <Online Large Model (Text Interaction)>.

1. **Prepare the video file:**

Place the video file to be tested in the following path:

```
~/yahboom_ws/src/largemodel/resources_file/analyze_video
```

Then name the video `test_video.mp4`

2. **Start the largemodel main program (text mode):**

Open a terminal and run the following command:

```
ros2 launch largemodel largemodel_control.launch.py text_chat_mode:=true
```

3. **Send a text command:**

Open another terminal and run the following command:

```
ros2 run text_chat text_chat
```

Then start typing: "Analyze this video."

4. **Observation:**

In the first terminal where the main program is run, you will see log output indicating that the system received the command, called the `analyze_video` tool, extracted keyframes, and finally printed out the AI's summary of the video content.

4. Common Problems and Solutions

Issue 1: Error message "Cannot find video file" or "Unable to extract keyframes from video."

Solution:

1. **Check Path:** Ensure you have placed the video file in the specified path, named it `test_video.mp4`, and have read permissions.
2. **Video Format/Encoding:** Ensure the video file is not corrupted and is in .mp4 format.

Issue 2: Analyzing a long video is very time-consuming.

Solution:

1. **Use a lighter model:** If hardware resources are limited, switch to a video understanding model with fewer parameters.