

Multimodal Autonomous Agent Application

Multimodal Autonomous Agent Application

1. Concept Introduction
 - 1.1 What is an "Autonomous Agent"?
 - 1.2 Implementation Principles
2. Code Analysis
 - Key Code
 1. Agent Core Workflow (`largetmodel/utils/ai_agent.py`)
 2. Task Planning and LLM Interaction (`largetmodel/utils/ai_agent.py`)
 3. Parameter Processing and Data Flow Implementation (`largetmodel/utils/ai_agent.py`)
 - Code Analysis
3. Practical Operations
 - 3.1 Configuring the Offline Large Model
 - 3.1.1 Configuring the LLM Platform (`yahboom.yaml`)
 - 3.1.2 Configuring the Model Interface (`larget_model_interface.yaml`)
 - 3.2 Launching and Testing the Function (Text Input Mode)
4. Common Problems and Solutions
 - 4.1 Abnormal Agent Behavior
 - Problem 1: The agent is stuck in an infinite loop or repeatedly executing the same tool.

Note: Due to performance limitations, the Raspberry Pi 5 1GB, 2GB, and 4GB versions cannot run offline. You can refer to the tutorials for online large models.

1. Concept Introduction

1.1 What is an "Autonomous Agent"?

In the `largetmodel` project, **multimodal autonomous agents** represent the most advanced form of intelligence. They no longer simply respond to a user's command; instead, they are capable of **autonomously thinking, planning, and continuously invoking multiple tools to achieve a complex goal**.

The core of this functionality is the `agent_call` tool or its **underlying** `ToolChainManager`. When a user issues a complex request that cannot be accomplished with a single tool call, the autonomous agent is activated.

1.2 Implementation Principles

The autonomous agent implementation in `largetmodel` follows the industry-leading **ReAct (Reason + Act)** paradigm. Its core idea is to mimic the human problem-solving process, cycling between "reflection" and "action."

1. **Reason:** When the agent receives a complex goal, it first invokes a powerful language model (LLM) to "reflect." It asks itself, "What should I do first to achieve this goal? Which tool should I use?" The LLM's output isn't a final answer, but rather an action plan.
2. **Act:** Based on the LLM's reflection, the agent performs the corresponding action—calling the `ToolsManager` to run a specified tool (such as `visual_positioning`).

3. **Observe:** The agent obtains the result of the previous action (the "observation"), for example, `{"result": "The cup was found at [120, 300, 180, 360]"}`.
4. **Reflect:** The agent submits this observation, along with the original goal, to the LLM for a second round of reflection. It asks itself, "I've found the location of the cup. What should I do next to learn its color?" The LLM might generate a new action plan, such as `{"thought": "I need to analyze the image of the area where the cup is located to determine its color", "action": "seewhat", "args": {"crop_area": [120, 300, 180, 360]}}`.

This **think -> act -> observe** cycle continues until the initial goal is achieved, at which point the agent generates and outputs the final answer.

2. Code Analysis

Key Code

1. Agent Core Workflow (`largemodel/utils/ai_agent.py`)

The `_execute_agent_workflow` function is the agent's main execution loop, defining the core "plan -> execute" process.

```
# From largemodel/utils/ai_agent.py

class AIAgent:
    # ...

    def _execute_agent_workflow(self, task_description: str) -> Dict[str, Any]:
        """
        Executes the agent workflow: Plan -> Execute.
        """
        try:
            # Step 1: Mission Planning
            self.node.get_logger().info("AI Agent starting task planning phase")
            plan_result = self._plan_task(task_description)

            # ... (Return early if planning fails)

            self.task_steps = plan_result["steps"]

            # Step 2: Follow all steps in order
            execution_results = []
            tool_outputs = []

            for i, step in enumerate(self.task_steps):
                # 2.1. Process data references in parameters before execution
                processed_parameters =
self._process_step_parameters(step.get("parameters", {}), tool_outputs)
                step["parameters"] = processed_parameters

                # 2.2. Execute a single step
                step_result = self._execute_step(step, tool_outputs)
                execution_results.append(step_result)

                # 2.3. If the step succeeds, save its output for reference in
                subsequent steps
```

```

        if step_result.get("success") and step_result.get("tool_output"):
            tool_outputs.append(step_result["tool_output"])
        else:
            # If any step fails, abort the entire task
            return { "success": False, "message": f"Task terminated
because step '{step['description']}' failed." }

        # ... Summarize and return the final result
        summary = self._summarize_execution(task_description,
execution_results)
        return { "success": True, "message": summary, "results":
execution_results }

    # ... (Exception handling)

```

2. Task Planning and LLM Interaction (`largemodel/utils/ai_agent.py`)

The core of the `_plan_task` function is to build a sophisticated prompt, leveraging the large model's inherent reasoning capabilities to generate a structured execution plan.

```

# From largemodel/utils/ai_agent.py

class AIAgent:
    # ...
    def _plan_task(self, task_description: str) -> Dict[str, Any]:
        """
        Uses the large model for task planning and decomposition.
        """

        # Dynamically generate a list of available tools and their descriptions
        tool_descriptions = []
        for name, adapter in self.tools_manager.tool_chain_manager.tools.items():
            # ... (Get tool description from adapter.input_schema)
            tool_descriptions.append(f"- {name}({params}): {description}")
        available_tools_str = "\n".join(tool_descriptions)

        # Build a highly structured plan
        planning_prompt = f"""
As a professional task planning agent, please break down user tasks into a series
of specific, executable JSON steps.

## Available Tools:##
{available_tools_str}

## Core Rules:##
1. ##Data Passing##: When a subsequent step requires the output of a previous
step, it must be referenced using the `{{{steps.N.outputs.KEY}}}` format.
- `N` is the step ID (starting at 1).
- `KEY` is the specific field name in the output data of the previous step.
2. ##JSON Format##: The return value must be a strict JSON object.

## User Task:##
{task_description}
"""

```

```

# Call the large model for planning
messages_to_use = [{"role": "user", "content": planning_prompt}]
# Note that this calls the general text inference API.
result = self.node.model_client.infer_with_text("",
message=messages_to_use)

# ... (Parse the JSON response and return a list of steps)

```

3. Parameter Processing and Data Flow Implementation (`largemodel/utils/ai_agent.py`)

The `_process_step_parameters` function is responsible for parsing placeholders and implementing data flow between steps.

```

# From largemodel/utils/ai_agent.py

class AIAgent:
    # ...
    def _process_step_parameters(self, parameters: Dict[str, Any],
previous_outputs: List[Any]) -> Dict[str, Any]:
    """
    Parses parameter dictionary, finds and replaces all {{...}} references.
    """
    processed_params = parameters.copy()
    # Regular expression used to match placeholders in the format
    {{steps.N.outputs.KEY}}
    pattern = re.compile(r"\\{\\{steps\\.(\\d+)\\.outputs\\.(.+?)\\}\\}\\}")

    for key, value in processed_params.items():
        if isinstance(value, str) and pattern.search(value):
            # Use re.sub and a replacement function to process all found
            placeholders
            # The replacement function will look up and return the value from
            the previous_outputs list
            processed_params[key] = pattern.sub(replacer_function, value)

    return processed_params

```

Code Analysis

The AI Agent is the "brain" of the system, translating high-level, sometimes ambiguous, tasks posed by the user into a precise, ordered series of tool calls. Its implementation is independent of any specific model platform and built on a general, extensible architecture.

1. **Dynamic Task Planning:** The Agent's core capability lies in the `_plan_task` function. Rather than relying on hard-coded logic, it dynamically generates task plans by interacting with a larger model.
 - **Self-Awareness and Prompt Construction:** At the beginning of planning, the Agent first examines all available tools and their descriptions. It then packages this tool information, the user's task, and strict rules (such as data transfer format) into a highly structured `planning_prompt`.

- **Model as Planner:** This prompt is fed into a general text-based model. The model reasoned based on the provided context and returned a multi-step action plan in JSON format. This design is highly scalable: as tools are added or modified in the system, the Agent's planning capabilities are automatically updated without requiring code modifications.
2. **Toolchain and Data Flow:** Real-world tasks often require the collaboration of multiple tools. For example, "take a picture and describe" requires the output (image path) of the "take a picture" tool to be used as the input of the "describe" tool. The AI Agent elegantly implements this through the `_process_step_parameters` function.
- **Data Reference Placeholders:** During the planning phase, large models embed special placeholders, such as `{{steps.1.outputs.data}}`, in parameter values where data needs to be passed.
 - **Real-Time Parameter Replacement:** In the `_execute_agent_workflow` main loop, `_process_step_parameters` is called before each step. It uses regular expressions to scan all parameters of the current step. Upon discovering a placeholder, it finds the corresponding data from the output list of the previous step and replaces it in real time. This mechanism is key to automating complex tasks.
3. **Supervised Execution and Fault Tolerance:** `_execute_agent_workflow` constitutes the Agent's main execution loop. It strictly follows the planned sequence of steps, executing each action sequentially and ensuring data is correctly passed between them.
- **Atomic Steps:** Each step is treated as an independent "atomic operation." If any step fails, the entire task chain immediately aborts and reports an error. This ensures system stability and predictability, preventing continued execution in an erroneous state.

In summary, the general implementation of the AI Agent demonstrates an advanced software architecture: rather than solving a problem directly, it builds a framework that enables an external, general-purpose reasoning engine (a large model) to solve the problem. Through two core mechanisms, dynamic programming and data flow management, the Agent orchestrates a series of independent tools into complex workflows capable of completing advanced tasks.

3. Practical Operations

3.1 Configuring the Offline Large Model

3.1.1 Configuring the LLM Platform (`yahboom.yaml`)

This file determines which large model platform the `model_service` node loads as its primary language model.

First, enter the command in the terminal to enter the Docker container:

```
./ros2_docker.sh
```

If you need to enter other commands in the same Docker container later, simply enter `./ros2_docker.sh` again in the host terminal.

1. **Open the file in the terminal:**

```
vim ~/yahboom_ws/src/largemodel/config/yahboom.yaml
```

2. Modify/Confirm `llm_platform`:

```
model_service:                                #Model server node parameters
  ros__parameters:
    language: 'zh'                             #Large Model Interface Language
    useolinetts: True                          #This item is invalid in text mode
    and can be ignored

    # Large model configuration
    llm_platform: 'ollama'                     # Key: Make sure it's 'ollama'
    regional_setting : "China"
    camera_type: 'csi'                         # Camera type: 'csi', 'usb'
    mjpeg_stream_url: 'http://172.17.0.1:8080/camera.mjpg' # MJPEG stream URL
    for CSI camera in Docker
```

3.1.2 Configuring the Model Interface (`large_model_interface.yaml`)

This file defines which vision model to use when the platform is selected as `ollama`.

1. Open the file in a terminal.

```
vim ~/yahboom_ws/src/largemodel/config/large_model_interface.yaml
```

2. Find the ollama-related configuration.

```
#.....
## Offline Large Language Models
# ollama Configuration
ollama_host: "http://localhost:11434" # ollama server address
ollama_model: "llava" # Important: Change this to the multimodal model you
downloaded, such as "llava"
#.....
```

Note: Please ensure that the model specified in the configuration parameters (e.g., `llava`) can handle multimodal input.

3.2 Launching and Testing the Function (Text Input Mode)

Note: The performance will be poor when using a model with a small number of parameters or when running at a very low memory usage. For a better experience with this function, please refer to the corresponding section in <Online Large Model (Text Interaction)>.

1. Launch the `largemodel` Main Program (Text Mode):

Open a terminal and run the following command:

```
ros2 launch largemodel largemodel_control.launch.py --text_chat_mode:=true
```

2. Send a Text Command:

Open another terminal and run the following command:

```
ros2 run text_chat text_chat
```

Then start typing: "Based on the current environment, save the generated environment description as a txt document."

3. **Observation Results:**

In the first terminal running the main program, you will see log output indicating that the system receives the text command, invokes the `aiagent` tool, and then provides a prompt to the LLM. The LLM will analyze the detailed steps of the tool invocation. For example, in this question, the `seewhat` tool will be called to obtain the image, which will then be parsed by the LLM. The parsed text will be saved in the

`~/yahboom_ws/src/largemodel/resources_file/documents` folder.

4. Common Problems and Solutions

4.1 Abnormal Agent Behavior

Problem 1: The agent is stuck in an infinite loop or repeatedly executing the same tool.

Solution:

1. **Enhance the Prompt:** Add stronger restrictions to the prompt, such as "Do not perform the same operation repeatedly" or "If the observation results are the same twice in a row, try a different tool or end the task."
2. **Replace a More Powerful LLM:** The agent's logical capabilities largely depend on the intelligence of the backend LLM. A more powerful model can better understand the task and make plans.